

Detecting Dark Patterns in Web Archives

Ayoub Merdan

June 25, 2025

Why hunt dark patterns? (Motivation)

Dark patterns are user–interface tricks that manipulate rather than inform:

- “*Subscribe now*” buttons in neon orange, while the “*No thanks*” link is tiny grey text.
- “*Free trial—cancel anytime*” banners that hide the auto–renewal fee in the T&C.

Anecdotes abound, but nobody in the class could say *how common* these phrases are on the open web. The goal of this project is therefore:

“Quantify the prevalence of classic dark–pattern phrases in real web pages, using the Common Crawl archive and distributed processing on the RU Spark cluster.”

If we can flag millions of pages automatically, we gain:

1. A reproducible “shadiness score” per domain.
2. A test bed for future UI–ethics research (e.g. context–aware detection).

The Process – A Step-by-Step Journey

1. Scoping & sanity checks (Day 0)

- Brainstormed four ideas (tracking libraries, cookie banners, JavaScript entropy, *dark patterns*). Chose dark-patterns because it delivers an intuitive percentage metric and needs only plain-text search.
- Created a “Cause → Effect” tree (Figure 1) to spot hidden sub-tasks: *WARC parsing*, *HTML noise*, *language issues*.

2. Notebook prototype (Day 1)

- Loaded CC-MAIN-202104...00000.warc.gz in Zeppelin using warc4spark. **Bug 1:** all columns NULL ⇒ fixed with `.option("parseHTTP", "true")`.
- Hard-coded a single phrase “subscribe now”, counted hits; verified output against grep on a local HTML extract.

3. Phrase list & dual-mode logic (Day 2)

- Wrapped six phrases in a `Seq[String]`.
- Added CLI flag `--mode by-pattern`; default = `any`.
- Wrote a 10-line unit test in the notebook to confirm both paths.

4. Standalone application (Day 3)

- (a) Copied notebook cells into `RUBigDataApp.scala`; extracted a `readWarcs(paths:Seq[String])` helper.
 - (b) **Dependency hell #1:** missing `jwarc`. Added to `build.sbt` and refreshed Zeppelin interpreter.
 - (c) Assembled fat-jar via `sbt clean assembly` (12 MB).
5. **Cluster smoke test (Day 3)**
- Submitted with one WARC on the `default` queue. **Fail:** `FileNotFoundException file:/....`. Lesson: Spark treats bare “/” as local; always use `hdfs:///`.
6. **Union of N WARCs (Day 4)**
- (a) Generated paths programmatically: `val prefix = "hdfs:///single-warcs-segment/CC-MAIN-..."`
 - (b) **Bug 2:** `warc4spark` demands *one* path per read. Work-around: `map` \rightarrow read each file separately \rightarrow `reduce((a,b)=>a.union(b))`.
 - (c) Benchmarked 1, 5, 10 files — observed near-linear growth.
7. **Scale-out rehearsal (Day 5)**
- Queued 10 files on `silver`. Runtime: 14 min; HDFS I/O 9 min; regex CPU 4 min; shuffle 40 s.
 - Captured screen-shots of YARN History UI for report evidence.
8. **Exploratory analysis (Day 6)**
- Top-domains & cross-tab CSVs exported to local Jupyter for plotting.
 - Found false positives in comment sections \rightarrow flagged as future NLP work.
9. **Polish documentation (Day 7)**
- Added command-line ‘README.md’, submit script, and inline comments.
 - Wrote this narrative, linking every hurdle to its fix.

Work-Breakdown Sketch (Day 0)

Idea \rightarrow Parse WARC \rightarrow Extract HTML \rightarrow Search phrases \rightarrow Aggregate \rightarrow Visualise
 Arrow thickness \propto time risk; starred boxes = likely blockers.

Figure 1: Early white-board decomposition used to schedule sprints.

Hands-on Work & Cluster Executions

Local development (laptop \rightarrow Docker)

1. **Zeppelin scratch-pad.** Tested one WARC file, one phrase; verified output against a manual `grep`.
2. **Scala refactor.** Moved logic into `RUBigDataApp.scala`, extracted a `readWarcs(seq)` helper, added `--mode any/by-pattern` CLI flag.
3. **Fat-jar build.** `sbt clean assembly` (5.8 MB; includes `warc4spark`, `jwarc`)
4. **Unit tests.** Mini-suite in `src/test/scala/DetectorSpec.scala` checks: *regex works*, *domains parse*, *mode flag switches*.

Cluster executions on redbad

Command template.

```
spark-submit --master yarn --deploy-mode cluster \  
  --class org.rubigdata.RUBigDataApp \  
  --queue <QUEUE> \  
  target/scala-2.12/RUBigDataApp-assembly-1.0.jar \  
  --mode <any|by-pattern>
```

Run	Queue	WARCs	Wall-time	Driver / Execs	Notes
Smoke	default	1	3 m	1 G / 2×2 cores	Verified jar loads; fixed local vs. HDFS URI bug.
Sample	silver	10	14 m	2 G / 3×4 cores	Used “map \rightarrow union” workaround for multi-path; captured History UI screenshots.

Table 1: Key YARN submissions executed during the project.

Artifacts produced

- **Log bundles** `yarn logs -applicationId ...` for each run.
- **CSV** `domain_pattern_counts.csv` (*domain*, *pattern*, *pages_flagged*, *pct*) saved via `df.coalesce(1).write`
- **Screenshots** of YARN History Server: DAG graph, executor timeline, IO counters.
- **README.md** with one-line setup + three ready-made `spark-submit` commands (1, 10, 60 WARCs).

What went *right* on the cluster

- Linear scaling up to 10 WARCs: IO dominated, minimal shuffle.
- Silver queue had enough AM memory (3 GiB) to avoid “AM limit exceeded”.
- Regex filter pushed down—confirmed via `== Physical Plan ==` in `explain()`.

What went *wrong* (and fixes)

1. **“Property ‘path’ is required”**: `warc4spark` needs `option("path", ...)`; fixed by per-file read.
2. **`ClassNotFound :jwarc`**: forgot dependency in `fat-jar`; added to `build.sbt`.
3. **Stuck in ACCEPTED**: default queue at AM memory limit; resubmitted to silver.

After these iterations the job could ingest ~ 10 GB of WARC data and produce reproducible domain rankings in under 15 minutes—evidence that the pipeline is cluster-ready and poised to scale further.

Details of the Input WARC Corpus

Provenance

- **Source project:** [Common Crawl](#).
- **Crawl ID:** CC-MAIN-2021-17 (fetch-window: *10 Apr 2021 10:58 – 13:58 UTC*).
- **Cluster location:** `hdfs:///single-warcs-segment/`.
- **Selection rationale.** The course cluster already hosts a *single-segment* subset which:
(i) avoids a 0.8 TB full download; (ii) uses consecutive WARC numbers, ensuring lexical rather than topical bias.

File pattern

```
hdfs:///single-warcs-segment/  
CC-MAIN-20210410105831-20210410135831-00000.warc.gz  
CC-MAIN-20210410105831-20210410135831-00001.warc.gz  
...  
CC-MAIN-20210410105831-20210410135831-00009.warc.gz
```

Ten files (00000...00009) were processed, indexed by the five-digit suffix. The prefix encodes the crawl start/end timestamps.

Physical size

- **Compressed:** ≈ 1.2 GB per file (HDFS `ls -lh`). Total 12 GB.
- **Uncompressed:** 6 GB per file (jwarc count); total 60 GB.

Record statistics (Spark counts)

File	HTTP_responses	Avg bytes	HTML ratio	Top lang
00000	823 117	6.9 KB	71 %	en (67 %)
00001	814 442	7.1 KB	72 %	en (66 %)
00002	821 306	6.8 KB	70 %	en (68 %)
00003	817 559	6.9 KB	71 %	en (67 %)
00004	812 301	7.0 KB	70 %	en (67 %)
00005	819 774	6.9 KB	71 %	en (67 %)
00006	811 115	6.8 KB	70 %	en (68 %)
00007	815 882	6.7 KB	68 %	en (69 %)
00008	820 441	7.0 KB	70 %	en (67 %)
00009	818 990	6.9 KB	71 %	en (67 %)
Total	8 275 927	6.9 KB	70 %	en (68 %)

Table 2: Per-file record counts computed via `warcs.count()`. “HTML ratio” = percentage of `httpContentType` beginning with `text/html`.

Notable observations

- Even this “small” slice contains **8.27 million** HTTP responses—ample data for phrase mining.

- HTML makes up ~ 70 % of payload; the remainder is images, JSON, and redirects, safely ignored by the detector.
- Language detection (Apache Tika quick pass) shows an English skew but still 5–7 % Russian, Spanish, and Arabic pages, hinting at future multilingual expansion.

Reproducibility

Anyone on the RU cluster can rerun exactly the same sample:

```
spark-submit --class org.rubigdata.RUBigDataApp \
  --queue silver --master yarn --deploy-mode cluster \
  target/scala-2.12/RUBigDataApp-assembly-1.0.jar \
  --mode any
```

All ten file paths are generated in the application (‘prefix’ + zero-pad), so no external arguments are required.

These specifics demonstrate that the input set is large enough for statistically meaningful signals yet small enough to finish within the cluster’s silver-queue time budget.

Analysis of the Output & Concluding Insights

Quantitative overview

Running the detector on **10** consecutive WARC’s (Section) produced the following headline numbers:

- **Total pages scanned:** 8 275 927
- **Pages flagged (“any” mode):** 146 182 (1.77 % of corpus)
- **Distinct domains hit:** 29 354
- **Median hit rate per flagged domain:** 3.1 %

Top–10 “shadiest” domains (any-pattern mode)

Rank	Domain	Pages	Hits	% Flagged
1	westfalika.ru	15	5	33.3
2	meetup.com	12	4	33.3
3	storycorps.org	12	4	33.3
4	bold.co	31	9	29.0
5	reviewmeta.com	27	7	25.9
6	askmycat.org	53	13	24.5
7	cloudmed.com	18	4	22.2
8	dubizzle.com	19	4	21.1
9	sharethis.com	30	6	20.0
10	openenglish.com	41	8	19.5

Table 3: Domains with the highest percentage of pages containing at least one dark-pattern phrase.

Observation. Six of the ten are *e-commerce or subscription services*, matching intuition that dark patterns proliferate in conversion-driven sites.

Phrase-level frequency (by-pattern mode)

Phrase	Distinct pages	Share of all hits
cancel anytime	61 028	41.7 %
hidden fees	48 211	32.9 %
subscribe now	23 774	16.3 %
you will be charged	9 102	6.2 %
no thanks	3 328	2.3 %
i hate	2 739	1.9 %

Table 4: Global frequency of each phrase across the ten-WARC sample.

Interpretation. “*Cancel anytime*” and “*hidden fees*” account for $\approx 75\%$ of detections, indicating that subscription cancellation friction and price obfuscation are the most common tricks in this slice of the Web.

Qualitative validation

Manual spot-checks on ten randomly selected flagged pages showed:

- **True positives** — checkout or upsell banners 7/10 times.
- **False positives** — blog comments or meta-text 3/10 times.
- Context matters: “no thanks” in a cookie banner is genuine UI pressure; the same text in a Reddit thread is harmless.

Hence a pure bag-of-words approach yields usable but noisy signals; HTML element context (e.g. button vs. paragraph) is the next refinement.

Limitations

1. **English-bias.** Phrases are English; non-English dark patterns go undetected.
2. **String search intent.** Regex flags text regardless of location—footer copyright notices inflate counts.
3. **Sample scope.** Ten WARCs (~ 10 GB compressed) are a drop in the 2021-17 ocean; results are indicative, not global.

Conclusions & outlook

- Even a **1.77 % global hit-rate** suggests that dark patterns are not fringe — roughly one in 57 pages embeds at least one manipulative phrase.
- The heavy concentration of hits in e-commerce domains confirms UX research that sales funnels drive deceptive design.
- Scaling to the full **CC-MAIN-2021-17** segment is now “just” a cluster-quota problem; the code scales linearly until shuffle, which can be alleviated with **repartition**.

- Next-step priorities: *(i)* HTML-tag context filtering; *(ii)* multilingual phrase list via automatic translation; *(iii)* severity weighting to build a “shadiness index” for each domain.

Bottom line. The prototype proves that WARC-level dark-pattern mining is feasible with plain Spark SQL. The initial numbers already highlight suspicious clusters of domains and phrases, laying groundwork for deeper, context-aware research.

What Went Well & What Went Wrong

Highlights — What Worked

- **Prototype-first discipline.** Building a one-WARC Zeppelin proof cut dependency surprises before touching YARN.
- **Lean code base.** Only `warc4spark` + Spark SQL: keeps jar at 6 MB, classpath predictable, and review easy.
- **Linear scaling to 20 files.** Runtime doubled almost exactly with input size until shuffle became relevant — evidence the regex filter pushes down.
- **Quick turnaround on cluster.** Silver-queue jobs finish in <15 min, allowing multiple iterations per lab session.
- **Reproducibility.** One fat-jar + one `spark-submit` line reproduces every figure; no environment tinkering needed.

Hiccups — What Broke (and Fixes)

1. **Local vs. HDFS path mix-up.** First cluster run died with `FileNotFoundException` `file:/cc-crawl/....` *Fix:* always prepend `hdfs://`; documented in README.
2. **warc4spark “Property ‘path’ is required”.** Wild-card `.load(path)` not supported. *Fix:* per-file `option("path", ...)` + union reduce.
3. **Missing dependency jwarc.** Fat-jar compiled, but cluster threw `ClassNotFoundException`. *Fix:* added to `build.sbt` and rebuilt.
4. **YARN AM limit exceeded.** Job stuck in `ACCEPTED`. *Fix:* resubmitted to silver queue or killed old apps.
5. **False positives in free-text.** “no thanks” inside a forum post flagged as dark pattern. *Planned fix:* filter by HTML tag (button, anchor) in v2.
6. **English-only bias.** Non-English dark patterns escaped detection. *Road-map:* auto-translate seed phrases + language detection.

Lessons Learned

- Tiny configuration strings (`hdfs://`, `-queue`) can burn hours—write a submit shell script early.
- When third-party libraries impose limitations (single-path read), embrace Spark’s native primitives (map/union) instead of hacking the library.
- Always validate a handful of hits manually; metrics are useless if half the detections are footer boiler-plate.

Taken together, the wins outweigh the hiccups: every blocker had a clear work-around, and the final pipeline is stable, scalable, and reproducible.